SDC2026 Vortex RunPod POU

Proof of RunPod Utilisation

Purpose: GPU serverless OCR extraction proof-of-use
Endpoint type: Queue-based RunPod Serverless
Invocation: authenticated POST to https://api.runpod.ai/v2/ojf4e2o9yj00cy/run or /runsync

Team Vortex · Vortex OrderGuard

Prepared 24 October 2026

Verified RunPod Evidence

A RunPod Serverless request completed successfully during the hackathon test. The recorded result showed status COMPLETED, gpu_used=true and GPU name NVIDIA GeForce RTX 5090. Recorded delay time was approximately 9.33 seconds and execution time approximately 5.57 seconds.

The worker is used only for OCR/extraction. OrderGuard performs deterministic comparison and keeps approval/rejection under human control. In the captured test, OCR execution succeeded but the structured invoice fields were empty; the project therefore treats OCR as optional input and retains manual confirmation/fallback rather than claiming perfect invoice recognition.

- RunPod Serverless endpoint deployed
- EasyOCR-based GPU inference
- Sample synthetic invoice stored at docs/sample-invoice.png
- No API key committed to GitHub
- OCR does not approve or reject supplier cases

Sample Output Schema

supplier_name · item · quantity · unit · unit_price · invoice_total

Supporting repository note: docs/RUNPOD_PROOF_OF_USE.md